

# Lab 9

Bingying Liu

3/30/2020

## 1. Linear Regression

### Frequentist Approach

```
data("clouds", package = "HSAUR3")
#head(clouds)

# frequentist approach
ols <- lm(rainfall ~ seeding * (sne + cloudcover + prewetness + echomotion) + time,
         data = clouds)
summary(ols)
```

```
##
## Call:
## lm(formula = rainfall ~ seeding * (sne + cloudcover + prewetness +
##     echomotion) + time, data = clouds)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5259 -1.1486 -0.2704  1.0401  4.3913
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -0.34624    2.78773   -0.124  0.90306
## seedingyes     15.68293    4.44627    3.527  0.00372 **
## sne            0.41981    0.84453    0.497  0.62742
## cloudcover     0.38786    0.21786    1.780  0.09839 .
## prewetness     4.10834    3.60101    1.141  0.27450
## echomotionstationary 3.15281    1.93253    1.631  0.12677
## time          -0.04497    0.02505   -1.795  0.09590 .
## seedingyes:sne  -3.19719    1.26707   -2.523  0.02545 *
## seedingyes:cloudcover -0.48625    0.24106   -2.017  0.06482 .
## seedingyes:prewetness -2.55707    4.48090   -0.571  0.57796
## seedingyes:echomotionstationary -0.56222    2.64430   -0.213  0.83492
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.205 on 13 degrees of freedom
## Multiple R-squared:  0.7158, Adjusted R-squared:  0.4972
## F-statistic: 3.274 on 10 and 13 DF, p-value: 0.02431
```

### Exercise 1: Interpret the significant coefficients (at the 0.05 significance level).

- Using 0.05 as significance level, we can see that “seedingyes” and “seedingyes:sne” are statistically significant.
- seedingyes: when sne, cloudcover, prewetness and echomotion are zero, the change of seeding from no to yes has will increase the amount of rainfall by 15.68.
- seedingyes:sne : Holding all other variables the same, for cloud with no seeding, the effect of one unit of increase in sne is 0.4198; for cloud with seeding, the effect of one unit of increase in sne is 0.4198 - 3.1953 = -2.7755

### Bayesian Approach

```

beta0.prior <- cauchy()
beta.prior <- cauchy()

stan_glm <- stan_glm(data = clouds,
  formula = rainfall ~ seeding * (sne + cloudcover + prewetness + echomotion) + time,
  family = gaussian(),
  prior = beta.prior,
  prior_intercept = beta0.prior,
  refresh = 0,
  refresh = 0)

summary(stan_glm)

```

```

##
## Model Info:
## function:      stan_glm
## family:        gaussian [identity]
## formula:       rainfall ~ seeding * (sne + cloudcover + prewetness + echomotion) +
##               time
## algorithm:     sampling
## sample:        4000 (posterior sample size)
## priors:        see help('prior_summary')
## observations:  24
## predictors:    11
##
## Estimates:
##               mean    sd  10%   50%   90%
## (Intercept)      1.4    2.9 -2.2   1.3   5.0
## seedingyes      11.7    4.7  5.6  11.9  17.7
## sne              0.0    0.9 -1.1   0.0   1.1
## cloudcover       0.3    0.2  0.0   0.3   0.6
## prewetness       3.9    3.6 -0.6   3.9   8.6
## echomotionstationary 2.7    1.9  0.3   2.7   5.0
## time            0.0    0.0 -0.1   0.0   0.0
## seedingyes:sne   -2.2    1.3 -3.9  -2.2  -0.5
## seedingyes:cloudcover -0.4    0.2 -0.7  -0.4  -0.1
## seedingyes:prewetness -2.8    4.4 -8.4  -2.8   2.9
## seedingyes:echomotionstationary -0.2    2.5 -3.4  -0.2   3.1
## sigma           2.4    0.5  1.8   2.3   3.0
##
## Fit Diagnostics:
##               mean    sd  10%   50%   90%
## mean_PPD 4.4    0.7  3.6   4.4   5.3
##
## The mean_ppd is the sample average posterior predictive distribution of the outcome variable (for details see help('summary.stanreg')).
##
## MCMC diagnostics
##               mcse Rhat n_eff
## (Intercept)    0.1  1.0  2353
## seedingyes      0.1  1.0  1774
## sne            0.0  1.0  2015
## cloudcover      0.0  1.0  1923
## prewetness      0.1  1.0  1808
## echomotionstationary 0.0  1.0  1931
## time           0.0  1.0  2505
## seedingyes:sne  0.0  1.0  1799
## seedingyes:cloudcover 0.0  1.0  1569
## seedingyes:prewetness 0.1  1.0  1830
## seedingyes:echomotionstationary 0.1  1.0  2216
## sigma          0.0  1.0  1161
## mean_PPD       0.0  1.0  3840
## log-posterior  0.1  1.0   626
##
## For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence Rhat=1).

```

## Exercise 2: How do the estimated coefficients compare in this glm model to those from the model fit using lm?

- lm has larger value in positive coefficients and has smaller value in negative coefficients (which means it's more extreme and more overfitting to the data than stan.glm.)

## Exercise 3: How do the credible intervals and standard errors of the coefficients compare?

```
# 95% CI comparison
round(confint(ols, level=0.95),3)
```

```
##                2.5 % 97.5 %
## (Intercept)    -6.369  5.676
## seedingyes      6.077 25.289
## sne            -1.405  2.244
## cloudcover     -0.083  0.859
## prewetness     -3.671 11.888
## echomotionstationary -1.022  7.328
## time           -0.099  0.009
## seedingyes:sne  -5.935 -0.460
## seedingyes:cloudcover -1.007  0.035
## seedingyes:prewetness -12.237  7.123
## seedingyes:echomotionstationary -6.275  5.150
```

```
round(posterior_interval(stan.glm, prob = 0.95), 3)
```

```
##                2.5% 97.5%
## (Intercept)    -4.195  7.312
## seedingyes      2.355 20.648
## sne            -1.777  1.698
## cloudcover     -0.086  0.773
## prewetness     -3.311 11.154
## echomotionstationary -1.024  6.381
## time           -0.092  0.013
## seedingyes:sne  -4.730  0.393
## seedingyes:cloudcover -0.880  0.071
## seedingyes:prewetness -11.715  5.737
## seedingyes:echomotionstationary -5.076  4.866
## sigma          1.636  3.497
```

- In general, the standard errors of stan.glm is bigger than lm. However, credible intervals of stan.glm and lm are very similar.

## 2. Logistic Regression

```
seed <- 196
admissions <- read.csv("https://stats.idre.ucla.edu/stat/data/binary.csv")
## view the first few rows of the data
# head(admissions)
admissions$rank <- factor(admissions$rank)
admissions$admit <- factor(admissions$admit)
admissions$gre <- scale(admissions$gre)
p <- 5
n <- nrow(admissions)
```

### Frequentist Approach

```
freq.mod <- glm(admit ~. , data = admissions,
                family = binomial())
summary(freq.mod)
```

```
##
## Call:
## glm(formula = admit ~ ., family = binomial(), data = admissions)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6268  -0.8662  -0.6388   1.1490   2.0790
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.6592     1.1657  -2.281 0.022535 *
## gre           0.2616     0.1264   2.070 0.038465 *
## gpa           0.8040     0.3318   2.423 0.015388 *
## rank2        -0.6754     0.3165  -2.134 0.032829 *
## rank3        -1.3402     0.3453  -3.881 0.000104 ***
## rank4        -1.5515     0.4178  -3.713 0.000205 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 458.52  on 394  degrees of freedom
## AIC: 470.52
##
## Number of Fisher Scoring iterations: 4
```

#### Exercise 4: Interpret the significant coefficients (at the 0.05 significance level).

- gre: Holding all else constant, for individual who is in rank 1, increasing gre by 1 unit is expected to increase the log-odds of being admitted by 0.2616.
- gpa: Holding all else constant, for individual who is in rank 1, increasing gpa by 1 unit is expected to increase the log-odds of being admitted by 0.8040.
- rank2: Holding all else constant, for individual who switch from rank 1 to rank 2, log-odds of being admitted is decreased by 0.6754.
- rank3: Holding all else constant, for individual who switch from rank 1 to rank 3, log-odds of being admitted is decreased by 1.3402.
- rank4: Holding all else constant, for individual who switch from rank 1 to rank 5, log-odds of being admitted is decreased by 1.5515.

#### Weakly Informative Prior: Normal


```
post1 <- stan_glm(admit ~ ., data = admissions,
                  family = binomial(link = "logit"),
                  prior = normal(0,1), prior_intercept = normal(0,1),
                  seed = seed,
                  refresh = 0)
summary(post1)
```

```
##
## Model Info:
## function:      stan_glm
## family:        binomial [logit]
## formula:       admit ~ .
## algorithm:     sampling
## sample:        4000 (posterior sample size)
## priors:        see help('prior_summary')
## observations:  400
## predictors:    6
##
## Estimates:
##           mean   sd   10%   50%   90%
## (Intercept) -2.8   1.1 -4.3   -2.8  -1.4
## gre          0.3   0.1  0.1    0.3   0.4
## gpa          0.8   0.3  0.4    0.8   1.2
## rank2       -0.5   0.3 -0.8   -0.5  -0.1
## rank3       -1.1   0.3 -1.5   -1.1  -0.7
## rank4       -1.2   0.4 -1.7   -1.3  -0.8
##
## Fit Diagnostics:
##           mean   sd   10%   50%   90%
## mean_PPD 0.3    0.0  0.3   0.3   0.4
##
## The mean_ppd is the sample average posterior predictive distribution of the outcome variable (for details see help('summary.stanreg')).
##
## MCMC diagnostics
##           mcse Rhat n_eff
## (Intercept)  0.0  1.0  3128
## gre          0.0  1.0  2999
## gpa          0.0  1.0  3199
## rank2        0.0  1.0  1752
## rank3        0.0  1.0  1597
## rank4        0.0  1.0  2337
## mean_PPD     0.0  1.0  4222
## log-posterior 0.0  1.0  1844
##
## For each parameter, mcse is Monte Carlo standard error, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence Rhat=1).
```

### Exercise 5: What do our choice of priors say about our beliefs? How do we interpret these normal priors?

- Choosing  $N(0, 1)$  means we believe that model coefficients/intercept are as likely to be positive as they are to be negative but they are highly unlikely to be far from zero.

```
mcmc_areas(as.matrix(post1), prob = 0.95, prob_outer = 1)
```



```
round(coef(post1), 3)
```

```
## (Intercept)      gre      gpa      rank2      rank3      rank4
##      -2.832      0.266      0.790     -0.462     -1.084     -1.252
```

```
round(posterior_interval(post1, prob = 0.95), 3)
```

```
##           2.5% 97.5%
## (Intercept) -5.156 -0.602
## gre          0.026 0.515
## gpa          0.173 1.443
## rank2        -1.066 0.116
## rank3        -1.712 -0.487
## rank4        -2.014 -0.531
```

```
round(confint(freq.mod, level=0.95),3) # comparison with the glm model
```

```
## Waiting for profiling to be done...
```

```
##           2.5 % 97.5 %
## (Intercept) -4.978 -0.397
## gre          0.016 0.512
## gpa          0.160 1.464
## rank2        -1.301 -0.057
## rank3        -2.028 -0.670
## rank4        -2.400 -0.754
```

**Exercise 6:** How do the estimated coefficients compare in this model to those from the model fit using glm?

- The estimated coefficients of stan\_glm are very similar (or slightly pulled towards 0) to those of glm since the bayesian regression model is using a weakly informative prior and data outweighs the prior belief.

**Exercise 7:** How do the credible intervals and standard errors of the coefficients compare to the confidence intervals and standard errors from the model fit using glm?

- The credible intervals and standard errors of `stan_glm` are also very similar to those of `glm` because of the same reasons above (data outweighs the prior belief).

## Posterior predictive checks

```
(loo1 <- loo(post1, save_psis = TRUE))
```

```
##
## Computed from 4000 by 400 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo   -235.3  8.6
## p_loo       5.6   0.3
## looic      470.6 17.2
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```

```
post0 <- stan_glm(admit ~ 1, data = admissions,
  family = binomial(link = "logit"),
  prior = normal(0,1), prior_intercept = normal(0,1),
  seed = seed,
  refresh = 0)
(loo0 <- loo(post0, save_psis = T))
```

```
##
## Computed from 4000 by 400 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo   -394.4 0.0
## p_loo      117.3 0.0
## looic      788.9 0.0
## -----
## Monte Carlo SE of elpd_loo is 0.5.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```

## Exercise 8: Which model is better? Why?

- `post1` is better than `post0` since it has all covariates. This means that covariates contain useful information for predictions.

```
preds <- posterior_linpred(post1, transform=TRUE)
pred <- colMeans(preds)

# classification accuracy
pr <- as.integer(pred >= 0.5)
round(mean(xor(pr, as.integer(admissions$admit==0))), 3)
```

```
## [1] 0.705
```

## The Horseshoe Prior

```
p0 <- 2 # prior guess for the number of relevant variables
tau0 <- p0/(p-p0) * 1/sqrt(n) # recommended by Pilronen and Vehtari (2017)
hs_prior <- hs(df=1, global_df=1, global_scale=tau0)
post2 <- stan_glm(admit ~ ., data = admissions,
  family = binomial(link = "logit"),
  prior = hs_prior, prior_intercept = normal(0,1),
  seed = seed,
  refresh = 0)


round(coef(post2), 3)
```

```
## (Intercept)      gre      gpa      rank2      rank3      rank4
##      -2.472      0.242      0.644     -0.191     -0.883     -1.051
```

```
round(posterior_interval(post2, prob = 0.95), 3)
```

```
##          2.5% 97.5%
## (Intercept) -5.015 -0.189
## gre         -0.003 0.517
## gpa         -0.020 1.375
## rank2       -0.967 0.208
## rank3      -1.697 -0.003
## rank4      -2.035 -0.008
```

```
mcmc_areas(as.matrix(post2), prob = 0.95, prob_outer = 1)
```



Exercise 9: How does posterior inference for the coefficients compare to when we used the weakly informative Normal prior above?

- The posterior inference for the coefficients are more pulled towards 0 when using horseshow prior than coefficients using weakly informative normal prior. Since Horseshoe places higher prior density on 0.

Exercise 10: How do the two models compare in terms of predictive performance? Consider using the loo function as we have been doing.

```
(loo2 <- loo(post2, save_psis = T))
```



```
##
## Computed from 4000 by 400 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo    -238.1  8.2
## p_loo         7.3  0.4
## looic        476.1 16.4
## -----
## Monte Carlo SE of elpd_loo is 0.1.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```

```
rstanarm::compare_models(loo1, loo2)
```

```
## Warning: 'rstanarm::compare_models' is deprecated.
## Use 'loo_compare' instead.
## See help("Deprecated")
```

```
## Warning: 'loo::compare' is deprecated.
## Use 'loo_compare' instead.
## See help("Deprecated")
```

```
## Model formulas:
## : NULL
## : NULLelpd_diff      se
##    -2.8        1.0
```

- The model using weakly informative normal prior performs better than model using horseshoe prior since elpd\_diff is negative (favors first model). Also since this dataset has  $n \gg p$  (number of observations greater than number of parameters), meaning that it doesn't quite make sense to use horseshoe prior which tends to shrink all coefficients towards 0.